

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

In [2]: titanic_data=pd.read_csv('titanic_train.csv')

In [3]: len(titanic_data)

Out[3]: 891

In [4]: titanic_data.head()

Out[4]:
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	NaN	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C85	C
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	NaN	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	C123	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	NaN	S

```


In [5]: titanic_data.index

Out[5]: RangeIndex(start=0, stop=891, step=1)

In [6]: titanic_data.columns

Out[6]: Index(['PassengerId', 'Survived', 'Pclass', 'Name', 'Sex', 'Age', 'SibSp',
              'Parch', 'Ticket', 'Fare', 'Cabin', 'Embarked'],
              dtype='object')

In [7]: titanic_data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   PassengerId      891 non-null    int64
1   Survived         891 non-null    int64
2   Pclass           891 non-null    int64
3   Name             891 non-null    object
4   Sex              891 non-null    object
5   Age              714 non-null    float64
6   SibSp            891 non-null    int64
7   Parch            891 non-null    int64
8   Ticket           891 non-null    object
9   Fare             891 non-null    float64
10  Cabin            204 non-null    object
11  Embarked         889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB

In [8]: titanic_data.dtypes

Out[8]: PassengerId      int64
Survived      int64
Pclass      int64
Name      object
Sex      object
Age      float64
SibSp      int64
Parch      int64
Ticket      object
Fare      float64
Cabin      object
Embarked      object
dtype: object

In [9]: titanic_data.describe()
```

Out[9]:	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare
count	891.000000	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
std	257.353842	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
min	1.000000	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	223.500000	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
50%	446.000000	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
75%	668.500000	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

In [10]:


```
titanic_data.isna()
```

Out[10]:	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
0	False	False	False	False	False	False	False	False	False	False	True	False
1	False	False	False	False	False	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False	False	False	False	True	False
3	False	False	False	False	False	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False	False	False	False	True	False
...
886	False	False	False	False	False	False	False	False	False	False	True	False
887	False	False	False	False	False	False	False	False	False	False	False	False
888	False	False	False	False	False	True	False	False	False	False	True	False
889	False	False	False	False	False	False	False	False	False	False	False	False
890	False	False	False	False	False	False	False	False	False	False	True	False

891 rows × 12 columns

In [11]:


```
titanic_data.isna().sum()
```

Out[11]:


```

PassengerId      0
Survived         0
Pclass           0
Name             0
Sex              0
Age             177
SibSp            0
Parch            0
Ticket           0
Fare             0
Cabin           687
Embarked         2
dtype: int64


```

In [12]:


```
sns.heatmap(titanic_data.isna())
```

Out[12]:


```
<Axes: >
```



```
In [13]: (titanic_data['Age'].isna().sum()/len(titanic_data['Age']))*100
```


```
Out[13]: 19.865319865319865
```

```
In [14]: (titanic_data['Cabin'].isna().sum()/len(titanic_data['Cabin']))*100
```

```
Out[14]: 77.10437710437711
```

```
In [15]: sns.displot(x='Age',data=titanic_data)
```

```
Out[15]: <seaborn.axisgrid.FacetGrid at 0x2028649e4e0>
```




```
In [17]: titanic_data['Age'].isna().sum()
```

```
Out[17]: 0
```

```
In [18]: sns.heatmap(titanic_data.isna())
```

```
Out[18]: <Axes: >
```



```
In [19]: titanic_data.drop('Cabin',axis=1,inplace=True)
```

```
In [20]: titanic_data.head()
```

Out[20]:	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Embarked
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	S
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...)	female	38.0	1	0	PC 17599	71.2833	C
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	S
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	S
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	S

```
In [21]: titanic_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 11 columns):
#   Column      Non-Null Count  Dtype
---  -
0   PassengerId  891 non-null    int64
1   Survived     891 non-null    int64
2   Pclass       891 non-null    int64
3   Name         891 non-null    object
4   Sex          891 non-null    object
5   Age         891 non-null    float64
6   SibSp        891 non-null    int64
7   Parch        891 non-null    int64
8   Ticket       891 non-null    object
9   Fare         891 non-null    float64
10  Embarked     889 non-null    object
dtypes: float64(2), int64(5), object(4)
memory usage: 76.7+ KB
```

```
In [22]: titanic_data.dtypes
```

```
Out[22]: PassengerId    int64
Survived      int64
Pclass        int64
Name          object
Sex           object
Age          float64
SibSp         int64
Parch         int64
Ticket        object
Fare          float64
Embarked      object
dtype: object
```

```
In [23]: gender=pd.get_dummies(titanic_data['Sex'],drop_first=True)
```

```
In [24]: titanic_data['Gender']=gender
```

```
In [25]: titanic_data.head()
```

```
Out[25]:
```

	PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Embarked	Gender
0	1	0	3	Braund, Mr. Owen Harris	male	22.0	1	0	A/5 21171	7.2500	S	True
1	2	1	1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	0	PC 17599	71.2833	C	False
2	3	1	3	Heikkinen, Miss. Laina	female	26.0	0	0	STON/O2. 3101282	7.9250	S	False
3	4	1	1	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	0	113803	53.1000	S	False
4	5	0	3	Allen, Mr. William Henry	male	35.0	0	0	373450	8.0500	S	True

```
In [26]: titanic_data.drop(['Name','Sex','Ticket','Embarked'],axis=1,inplace=True)
```

```
In [27]: titanic_data.head()
```

```
Out[27]:
```

	PassengerId	Survived	Pclass	Age	SibSp	Parch	Fare	Gender
0	1	0	3	22.0	1	0	7.2500	True
1	2	1	1	38.0	1	0	71.2833	False
2	3	1	3	26.0	0	0	7.9250	False
3	4	1	1	35.0	1	0	53.1000	False
4	5	0	3	35.0	0	0	8.0500	True

```
In [28]: x=titanic_data[['PassengerId','Pclass','Age','SibSp','Parch','Fare','Gender']]
y=titanic_data['Survived']
```

```
In [29]: y
```

```
Out[29]: 0      0
         1      1
         2      1
         3      1
         4      0
         ..
        886     0
        887     1
        888     0
        889     1
        890     0
        Name: Survived, Length: 891, dtype: int64
```

```
In [31]: titanic_data.isnull().sum()
```

```
Out[31]: PassengerId    0
         Survived       0
         Pclass        0
         Age           0
         SibSp         0
         Parch         0
         Fare          0
         Gender        0
         dtype: int64
```

```
In [32]: from sklearn.linear_model import LogisticRegression
         from sklearn.preprocessing import StandardScaler
```

```
In [34]: features_matrix = titanic_data.iloc[:, 0:34]
```

```
In [35]: target_vector = titanic_data.iloc[:, -1]
```

```
In [37]: print('The Features Matrix Has %d Rows And %d Column(s)%(features_matrix.shape))
         print('The Target Matrix Has %d Rows And %d Column(s)%(np.array(target_vector).reshape(-1, 1).shape))
```

```
The Features Matrix Has 891 Rows And 8 Column(s)
The Target Matrix Has 891 Rows And 1 Column(s)
```

```
In [38]: features_matrix_standardized = StandardScaler().fit_transform(features_matrix)
```

```
In [60]: Logistic_Regression_Model = algorithm.fit(features_matrix_standardized, target_vector)
```

```
In [ ]: observation = [[1, 0, 0.99539, -0.05889, 0.8524299999999999, 0.02306,
                        0.8339799999999999, -0.37708, 1.0, 0.0376,
                        0.8524299999999999, -0.17755, 0.59755, -0.44945, 0.60536,
                        -0.38223, 0.8435600000000001, -0.38542, 0.58212, -0.32192,
                        0.56971, -0.29674, 0.36946, -0.47357, 0.56811, -0.51171,
                        0.4107800000000003, -0.4616800000000003, 0.21266, -0.3409,
                        0.42267, -0.54487, 0.18641, -0.453]]
```

```
In [40]: from sklearn.linear_model import LogisticRegression

         algorithm = LogisticRegression(penalty=None, dual=False, tol=1e-4, C=1.0, fit_intercept=True, l1_ratio=None)
```

```
In [43]: print("""The Model Says The Probability Of The Observation We Passed Belonging To Class ['g'] Is %s""")
```

```
The Model Says The Probability Of The Observation We Passed Belonging To Class ['g'] Is %s
```

```
In [48]: from sklearn.model_selection import train_test_split
```

```
In [49]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.33, random_state=42)
```

```
In [50]: from sklearn.linear_model import LogisticRegression
```

```
In [51]: lr=LogisticRegression()
```

```
In [52]: lr.fit(x_train,y_train)
```

```
C:\Users\lekha\AppData\Roaming\Python\Python312\site-packages\sklearn\linear_model\_logistic.py:469: Convergence
Warning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

```
Increase the number of iterations (max_iter) or scale the data as shown in:
```

```
https://scikit-learn.org/stable/modules/preprocessing.html
```

```
Please also refer to the documentation for alternative solver options:
```

```
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression
```

```
n_iter_i = _check_optimize_result(
```

```
Out[52]: ▾ LogisticRegression ⓘ ?  
LogisticRegression()
```

```
In [53]: predict=lr.predict(x_test)
```

```
In [54]: from sklearn.metrics import confusion_matrix
```

```
In [55]: pd.DataFrame(confusion_matrix(y_test,predict),columns=['Predicted No','Predicted Yes'],index=['Actual No','Actual Yes'])
```

```
Out[55]:
```

	Predicted No	Predicted Yes
Actual No	151	24
Actual Yes	38	82

```
In [56]: from sklearn.metrics import classification_report
```

```
In [57]: print(classification_report(y_test,predict))
```

	precision	recall	f1-score	support
0	0.80	0.86	0.83	175
1	0.77	0.68	0.73	120
accuracy			0.79	295
macro avg	0.79	0.77	0.78	295
weighted avg	0.79	0.79	0.79	295

```
In [ ]: x = titanic_data.drop('g',axis=1)  
y = titanic_data['g']
```

```
In [ ]: df['g'].value_counts()
```

```
In [64]: from sklearn.model_selection import train_test_split  
x_train, x_test, y_train, y_test = train_test_split(x, y, train_size=0.7, random_state=42)  
x_train.shape, x_test.shape
```

```
Out[64]: ((623, 7), (268, 7))
```

```
In [73]: from sklearn.ensemble import RandomForestClassifier  
rfc = RandomForestClassifier()  
rfc.fit(x_train,y_train)
```

```
Out[73]: ▾ RandomForestClassifier ⓘ ?  
RandomForestClassifier()
```

```
In [74]: rf = RandomForestClassifier()
```

```
In [75]: params = {'max_depth': [2,3,5,10,20],  
'min_samples_leaf': [5,10,20,50,100,200],  
'n_estimators': [10,25,30,50,100,200]}
```

```
In [76]: from sklearn.model_selection import GridSearchCV  
grid_search = GridSearchCV(estimator=rf,param_grid=params,cv = 2, scoring="accuracy")  
grid_search.fit(x_train, y_train)
```

```
Out[76]: ▸ GridSearchCV ⓘ ?  
▸ estimator: RandomForestClassifier  
▸ RandomForestClassifier ⓘ
```

```
In [77]: grid_search.best_score_
```

```
Out[77]: 0.8218113612004287
```

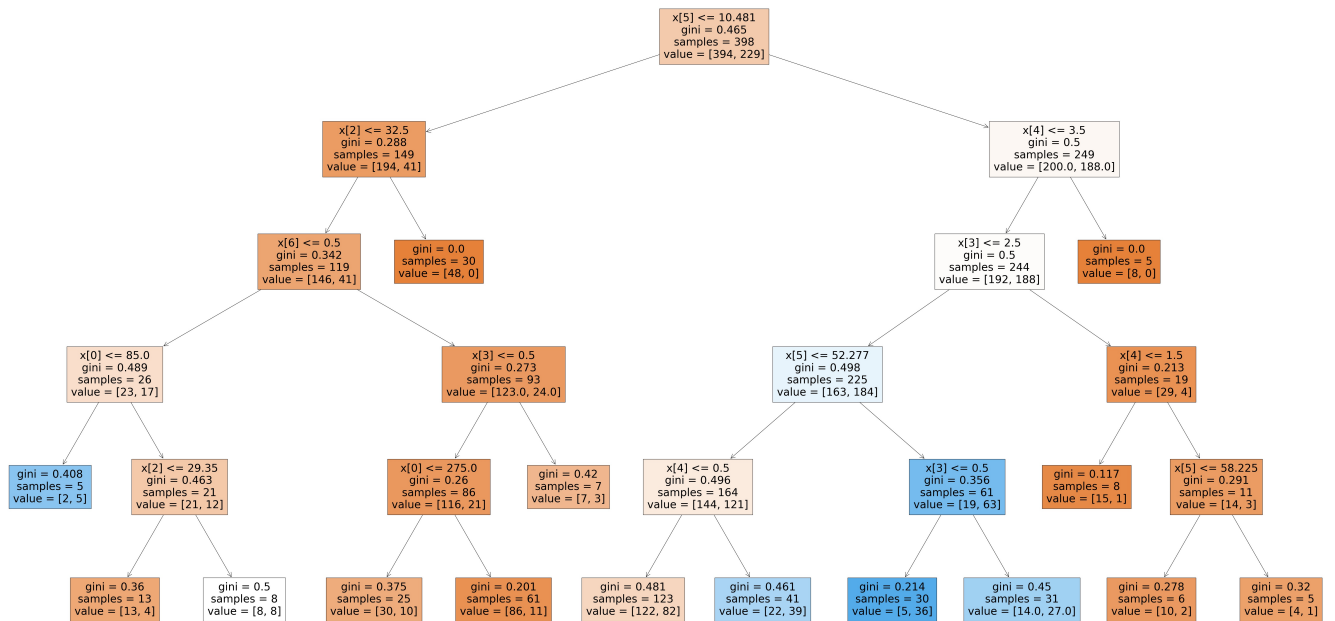
```
In [78]: rf_best = grid_search.best_estimator_  
print(rf_best)  
RandomForestClassifier(max_depth=5, min_samples_leaf=5, n_estimators=10)
```

```
In [79]: from sklearn.tree import plot_tree  
plt.figure(figsize=(80,40))
```



```
plot_tree(rf_best.estimators_[5], filled=True)
```

```
Out[79]: [Text(0.5357142857142857, 0.9166666666666666, 'x[5] <= 10.481\ngini = 0.465\nsamples = 398\nvalue = [394, 229]'),
),
Text(0.2857142857142857, 0.75, 'x[2] <= 32.5\ngini = 0.288\nsamples = 149\nvalue = [194, 41]'),
Text(0.23809523809523808, 0.5833333333333333, 'x[6] <= 0.5\ngini = 0.342\nsamples = 119\nvalue = [146, 41]'),
Text(0.09523809523809523, 0.4166666666666667, 'x[0] <= 85.0\ngini = 0.489\nsamples = 26\nvalue = [23, 17]'),
Text(0.047619047619047616, 0.25, 'gini = 0.408\nsamples = 5\nvalue = [2, 5]'),
Text(0.14285714285714285, 0.25, 'x[2] <= 29.35\ngini = 0.463\nsamples = 21\nvalue = [21, 12]'),
Text(0.09523809523809523, 0.08333333333333333, 'gini = 0.36\nsamples = 13\nvalue = [13, 4]'),
Text(0.19047619047619047, 0.08333333333333333, 'gini = 0.5\nsamples = 8\nvalue = [8, 8]'),
Text(0.38095238095238093, 0.4166666666666667, 'x[3] <= 0.5\ngini = 0.273\nsamples = 93\nvalue = [123.0, 24.0]'),
),
Text(0.3333333333333333, 0.25, 'x[0] <= 275.0\ngini = 0.26\nsamples = 86\nvalue = [116, 21]'),
Text(0.2857142857142857, 0.08333333333333333, 'gini = 0.375\nsamples = 25\nvalue = [30, 10]'),
Text(0.38095238095238093, 0.08333333333333333, 'gini = 0.201\nsamples = 61\nvalue = [86, 11]'),
Text(0.42857142857142855, 0.25, 'gini = 0.42\nsamples = 7\nvalue = [7, 3]'),
Text(0.3333333333333333, 0.5833333333333333, 'gini = 0.0\nsamples = 30\nvalue = [48, 0]'),
Text(0.7857142857142857, 0.75, 'x[4] <= 3.5\ngini = 0.5\nsamples = 249\nvalue = [200.0, 188.0]'),
Text(0.7380952380952381, 0.5833333333333333, 'x[3] <= 2.5\ngini = 0.5\nsamples = 244\nvalue = [192, 188]'),
Text(0.6190476190476191, 0.4166666666666667, 'x[5] <= 52.277\ngini = 0.498\nsamples = 225\nvalue = [163, 184]'),
),
Text(0.5238095238095238, 0.25, 'x[4] <= 0.5\ngini = 0.496\nsamples = 164\nvalue = [144, 121]'),
Text(0.47619047619047616, 0.08333333333333333, 'gini = 0.481\nsamples = 123\nvalue = [122, 82]'),
Text(0.5714285714285714, 0.08333333333333333, 'gini = 0.461\nsamples = 41\nvalue = [22, 39]'),
Text(0.7142857142857143, 0.25, 'x[3] <= 0.5\ngini = 0.356\nsamples = 61\nvalue = [19, 63]'),
Text(0.6666666666666666, 0.08333333333333333, 'gini = 0.214\nsamples = 30\nvalue = [5, 36]'),
Text(0.7619047619047619, 0.08333333333333333, 'gini = 0.45\nsamples = 31\nvalue = [14.0, 27.0]'),
Text(0.8571428571428571, 0.4166666666666667, 'x[4] <= 1.5\ngini = 0.213\nsamples = 19\nvalue = [29, 4]'),
Text(0.8095238095238095, 0.25, 'gini = 0.117\nsamples = 8\nvalue = [15, 1]'),
Text(0.9047619047619048, 0.25, 'x[5] <= 58.225\ngini = 0.291\nsamples = 11\nvalue = [14, 3]'),
Text(0.8571428571428571, 0.08333333333333333, 'gini = 0.278\nsamples = 6\nvalue = [10, 2]'),
Text(0.9523809523809523, 0.08333333333333333, 'gini = 0.32\nsamples = 5\nvalue = [4, 1]'),
Text(0.8333333333333333, 0.5833333333333333, 'gini = 0.0\nsamples = 5\nvalue = [8, 0]')]
```



```
In [80]: from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rf_best.estimators_[7],filled=True)
```


Out[82]:

	Varname	Imp
6	Gender	0.364141
5	Fare	0.229148
1	Pclass	0.112789
2	Age	0.111807
3	SibSp	0.076368
0	PassengerId	0.062734
4	Parch	0.043012

In []:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js